Lecture 12 Support Vector Machine II

STAT 441/505: Applied Statistical Methods in Data Mining

Linglong Kong

Department of Mathematical and Statistical Sciences
University of Alberta

Winter, 2016

UNIVERSITY OF ALBERTA
EDMONTON·ALBERTA·CANADA

Outline

Feature Expansion

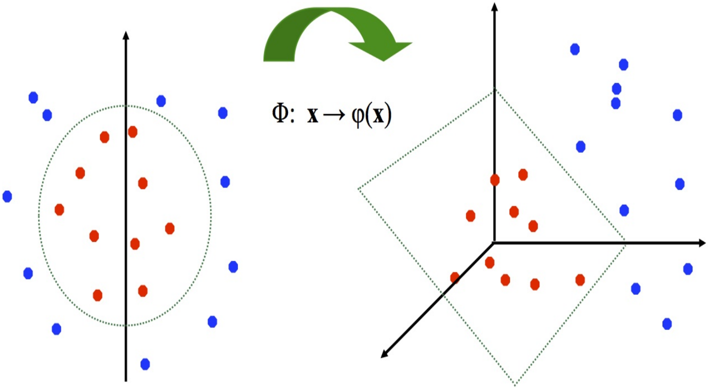

UNIVERSITY OF ALBERTA
EDMONTON·ALBERTA·CANADA

Feature Expansion

- Enlarge the space of features by including transformations; for example $X_1^2, X_2^3, X_1X_2, X_1X_2^2, \cdots$, Hence go from a p-dimensional space to a $M > p$ dimensional space.
- Fit a support-vector classifier in the enlarged space.
- This results in non-linear decision boundaries in the original space.
- Example: Suppose we use $(X_1, X_2, X_1^2, X_2^2, X_1X_2)$ instead of just (X_1, X_2). Then the decision boundary would be of the form

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 + \beta_4 X_2^2 + \beta_5 X_1 X_2 = 0.$$

- This leads to nonlinear decision boundaries in the original space (quadratic conic sections).

UNIVERSITY OF ALBERTA
EDMONTON·ALBERTA·CANADA

Cubic Polynomials

- Here we use a basis expansion of cubic polynomials — from 2 variables to 9.
- The support-vectorclassifier in the enlarged space solves the problem in the lower-dimensional space

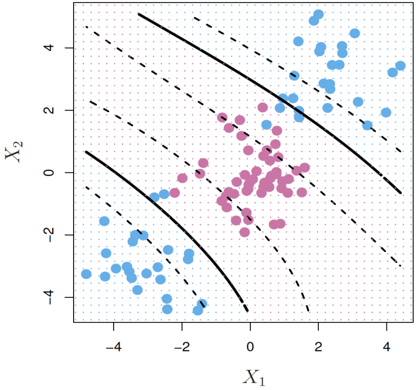

- The decision boundary is

$$\beta_0+\beta_1X_1+\beta_2X_2+\beta_3X_1^2+\beta_4X_2^2+\beta_5X_1X_2+\beta_6X_1^3+\beta_7X_2^3+\beta_8X_1X_2^2+\beta_9X_1^2X_2 = 0.$$

Nonlinearities and Kernels

- Polynomials (especially high-dimensional ones) get wild rather fast.
- There is a more elegant and controlled way to introduce nonlinearities in support vector classifier — through the use of kernels.
- Before we discuss these, we must understand the role of inner products in support vector classifier.

Inner products and kernels

- Inner product between vectors

$$\langle x_i, x_{i'} \rangle = \sum_j x_{ij} x_{i''j}.$$

- The linear support vector classifier can be represented as

$$f(x) = \beta_0 + \sum_i \alpha_i \langle x, x_i \rangle$$

- To estimate parameters $\alpha_1, \cdots, \alpha_n$ and β_0, all we need are $\binom{n}{2}$ inner products $\langle x, x_i \rangle$ between all pairs of training observations.
- It turns out that most of the $\hat{\alpha}_i$ can be zero

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \hat{\alpha}_i \langle x, x_i \rangle,$$

where $\mathcal{S}$ is the support set of indices i such that $\hat{\alpha}_i > 0$.

UNIVERSITY OF ALBERTA
EDMONTON·ALBERTA·CANADA

Kernels and Support Vector Machine

- If we can compute inner products between observations, we can fit a support vector classifier — can be very abstract!
- Some special kernel function can do this for us. E.g.

$$K(x_i, x_{i'}) = (1 + \sum_j x_{ij}x_{i"j})^2$$

computes the inner products needed for d dimensional polynomials — $\binom{p+d}{d}$ basis functions!

- The solotion has the form

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \hat{\alpha}_i K(x, x_i).$$

UNIVERSITY OF ALBERTA
EDMONTON·ALBERTA·CANADA

Radial Kernel

- The radial Kernel has the format

$$K(x_i, x_{i'}) = \exp\left(-\gamma \sum_j (x_{ij} - x_{i'j})^2\right),$$

where γ is tuning parameter.

- The decision bounady is,

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \hat{\alpha}_i \langle x, x_i \rangle,$$

implicit feature space; very high dimensional.

- Controls variance by squaring down most dimenions severely.

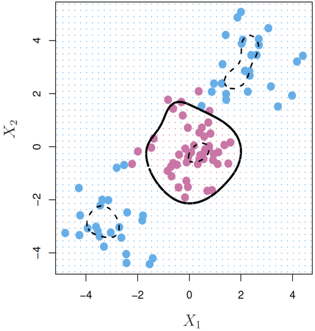

Example - Heart Data

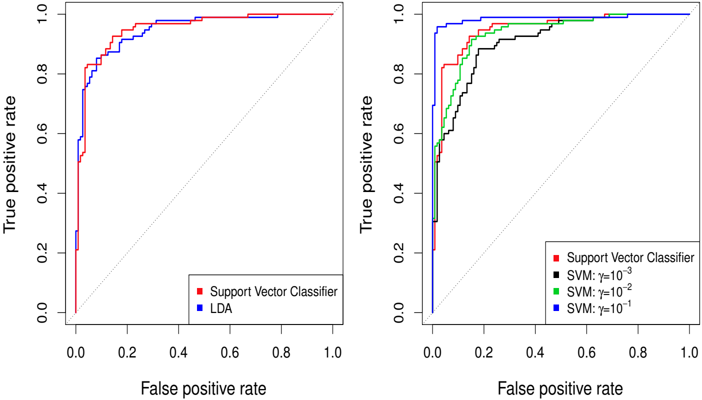

ROC curves on Training data

- ROC curve is obtained by changing the threshold 0 to threshold t in $\hat{f}(X) > t$, and recording false positive and true positive rates as t varies.

Example - Heart Data

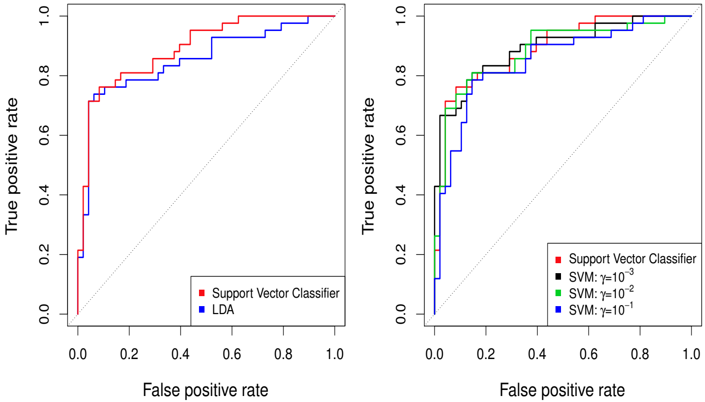

ROC curves on Testing data

UNIVERSITY OF ALBERTA
EDMONTON·ALBERTA·CANADA

SVMs: More than 2 classes

- The SVM as defined works for $K = 2$ classes. What do we do if we have $K > 2$ classes?
- OVA - One versus All. Fit K different 2-class SVM classifiers $\hat{f}_k(x),\ k = 1, \cdots, K$; each class versus the rest. Classify x^* to the class for which $\hat{f}_k(x^*)$ is largest.
- OVO - One versus One. Fit all $\binom{K}{2}$ pairwise classifiers $\hat{f}_{kl}(x)$. Classify x^* to the class that wins the most pairwise competitions.
- Which one to choose? If K is not too large, use OVO.

UNIVERSITY OF ALBERTA
EDMONTON·ALBERTA·CANADA

Support Vector Machine Versus Logistic Regression

- Let $f(X) = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p$, support vector machine can be rephrased as

$$\text{minimize}_{\beta_0,\beta_1,\cdots,\beta_p} \left\{ \sum_i \max[0, 1 - y_i f(x_i)] + \lambda \sum_j \beta_j^2 \right\},$$

where γ is tuning parameter.

- This has the form of loss plus penalty.
- The loss is known as hinge loss.
- Very similar to the loss in logistic regression (negative log-likelihood).

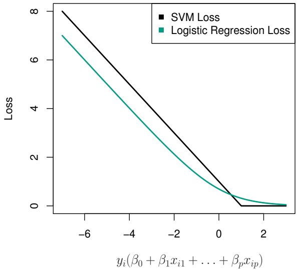

Kernels and Support Vector Machine

- When classes are (nearly) separable, SVM does better than LR. So does LDA.
- When not, LR (with ridge penalty) and SVM very similar.
- If you wish to estimate probabilities, LR is the choice.
- For nonlinear boundaries, kernel SVMs are popular. Can use kernels with LR and LDA as well, but computations are more expensive.

Summary and Remark

- Feature expansion
- Kernel Trick
- More than 2 classes
- Read textbook Chapter 12 and R code
- Do R lab