

# Lecture 22 Neural Network II

## STAT 441/505: Applied Statistical Methods in Data Mining

Linglong Kong

Department of Mathematical and Statistical Sciences  
University of Alberta

Winter, 2016

# Outline

Fitting Neural Networks

Training Neural Networks

Zip code data

Summary and Remark

# Fitting Neural Networks

- ▶ The unknown parameters in neural network model are called **weights**, denoted by  $\theta$ , which includes

$$\begin{aligned} &\{\alpha_{0m}, \alpha_m; m = 1, 2, \dots, M\} \quad M(p + 1) \text{ weights,} \\ &\{\beta_{0k}, \beta_k; k = 1, 2, \dots, K\} \quad K(M + 1) \text{ weights.} \end{aligned}$$

- ▶ In regression, we minimize RSS

$$R(\theta) = \sum_{i=1}^n R_i = \sum_{i=1}^n \sum_{k=1}^K (y_{ik} - f_k(x_i))^2.$$

- ▶ In classification, we minimize cross-entropy (deviance)

$$R(\theta) = \sum_{i=1}^n R_i = - \sum_{i=1}^n \sum_{k=1}^K f_k(x_i) \log f_k(x_i),$$

and the corresponding classifier is  $G(x) = \arg \max_k f_k(x)$ .

# Back-Propagation

- ▶ The generic approach to minimizing  $R(\theta)$  is by gradient descent, called **back-propagation** in this setting.
- ▶ Let  $z_{mi} = \sigma(\alpha_0 m + \alpha_m^T x_i)$  and  $z_i = (z_{1i}, \dots, z_{Mi})$ . The derivatives of  $R(\theta)$  are

$$\frac{\partial R_i}{\partial \beta_{km}} = -2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)z_{mi},$$

$$\frac{\partial R_i}{\partial \alpha_{kl}} = \sum_{k=1}^K 2y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)\beta_{km}\sigma'(\alpha_m^T x_i)x_{il}.$$

- ▶ which can be rewritten as

$$\frac{\partial R_i}{\partial \beta_{km}} = \delta_{ki}z_{mi}, \quad \frac{\partial R_i}{\partial \alpha_{kl}} = s_{mi}x_{il}, \quad (1.1)$$

where the quantities  $\delta_{ki}$  and  $s_{mi}$  are **errors** from the current model at the output and hidden layer units, respectively.

# Back-Propagation

- It can easily be shown that

$$s_{mi} = \sigma'(\alpha^T x_i) \sum_{k=1}^K \beta_{km} \delta_{ki}, \quad (1.2)$$

known as the **back-propagation equation**.

- Given the derivatives, a gradient decent update at  $(r + 1)$  iteration has the form

$$\beta_{km}^{(r+1)} = \beta_{km}^{(r)} - \gamma_r \sum_{i=1}^n \frac{\partial R_i}{\partial \beta_{km}^{(r)}}, \quad (1.3)$$

$$\alpha_{ml}^{(r+1)} = \alpha_{ml}^{(r)} - \gamma_r \sum_{i=1}^n \frac{\partial R_i}{\partial \alpha_{kl}^{(r)}},$$

where where  $\gamma_m$  is the **learning rate**.

# Back-Propagation

- ▶ In the **forward pass**, the current weights are fixed and the predicted values  $\hat{f}_k(x_i)$  are computed from formula in the last lecture.
- ▶ In the **backward pass**, the errors  $\delta_{ki}$  are computed, and then back-propagated via (1.2) to give the errors  $s_{mi}$ .
- ▶ Both sets of errors are then used to compute the gradients for the updates in (1.3) via (1.1).
- ▶ This two-pass procedure is what is known as **back-propagation** or **delta rule**.
- ▶ Back-propagation can be slow. Other methods include second-order techniques, conjugate gradients and variable metric methods.

# Training Neural Networks

- ▶ **Starting Values.** Usually starting values for weights are chosen to be random values near zero. Hence the model starts out nearly linear, and becomes nonlinear as the weights increase.
- ▶ **Overfitting.** Often neural networks have too many weights and will overfit the data at the global minimum of  $R$ .
- ▶ A more explicit method for regularization is **weight decay**, which is analogous to ridge regression, that is


$$J(\theta) = \sum_{km} \beta_{km}^2 + \sum_{ml} \alpha_{ml}^2.$$

- ▶ Other penalties are proposed as well, for example, the **weight elimination** penalty


$$J(\theta) = \sum_{km} \beta_{km}^2 / (1 + \beta_{km}^2) + \sum_{ml} \alpha_{ml}^2 / (1 + \alpha_{ml}^2).$$

# Training Neural Network

Neural Network - 10 Units, No Weight Decay



Neural Network - 10 Units, Weight Decay=0.02



The broken purple boundary is the Bayes error rate. Both use the softmax activation function and cross-entropy error.



# Training Neural Network

- ▶ **Number of hidden units and layers.** Generally speaking it is better to have too many hidden units than too few.
- ▶ **Multiple Minima.** The loss function  $R(\theta)$  is nonconvex and hence possesses many local minima.
- ▶ One must at least try a number of random starting configurations, and choose the solution giving lowest (penalized) error.
- ▶ Another approach is via **bagging**, which averages the predictions of networks training from randomly perturbed versions of the training data.


# Training Neural Network

- ▶ In summary, there are two free parameters to select: the weight decay  $\lambda$  and number of hidden units  $M$  as in

$$R(\theta) + \lambda J(\theta).$$


- ▶ As a learning strategy, one could fix either parameter at the value corresponding to the least constrained model, to ensure that the model is rich enough, and use cross-validation to choose the other parameter.

## Zip code data



Examples of training cases from ZIP code data. Each image is a  $16 \times 16$  8-bit grayscale representation of a handwritten digit.


# Zip code data



# Zip code data

- ▶ **Net-1:** No hidden layer, equivalent to multinomial logistic regression.
- ▶ **Net-2:** One hidden layer, 12 hidden units fully connected.
- ▶ **Net-3:** Two hidden layers locally connected.
- ▶ **Net-4:** Two hidden layers, locally connected with weight sharing.
- ▶ **Net-5:** Two hidden layers, locally connected, two levels of weight sharing.

# Zip code data



# Summary and Remark

- ▶ Back propagation
- ▶ Training neural network
- ▶ Zip code data
- ▶ Read textbook Chapter 11 and R code
- ▶ Do R lab